


Notes on Linear Regression

Luis Mendez

I. ASSUMPTIONS OF LINEAR REGRESSION

- Linearity: The relationship between the independent and dependent variable is linear.
- Independence: The residuals (errors) are independent of each other.
- Homoscedasticity: The residuals have constant variance at every level of the independent variable.
- Normality: The residuals are normally distributed.
- Multivariate Normality: The independent variables are normally distributed (for multiple regression).
- No Multicollinearity: The independent variables are not highly correlated with each other.



II. BUILDING A MODEL

- Define the problem and collect data.
- Preprocess the data (handle missing values, encode categorical variables, etc.).
- Split the data into training and testing sets.
- Train the linear regression model on the training data.
- Evaluate the model using metrics such as Mean Absolute Error (MAE), Mean Squared Error (MSE), or R-squared on the testing data.
- Interpret the coefficients to understand the relationship between independent and dependent variables.

BACKWARD ELIMINATION

- Step 1: Select a significance level (e.g., 0.05) to stay in the model.
- Step 2: Fit the model and calculate the p-values for each coefficient.
- Step 3: Remove the variable with the highest p-value (greater than a significance level, e.g., 0.05).
- Step 4: Repeat the process until all remaining variables have p-values below the significance level.
- Step 5: The final model will contain only the statistically significant variables. Fit the model without the variables that were removed and evaluate its performance.